



# Hatchet: Introduction

Abhinav Bhatele, University of Maryland



# Code-centric performance analysis


- Understanding performance bottlenecks is critical to optimizing (parallel) software
- Profiling and tracing tools help identify parts of code that consume the most time

gprof Data

4 bytes per bucket, each sample counts as 10.000ms

Name (location)	Samples	Calls	Time/Call	% Time
▼ Summary	2228			100.0%
► calc.c	590			26.48%
► copy.c	0			0.0%
► diag.c	25			1.12%
► main.c	0			0.0%
► time.c	653			29.31%
▼ tstep.c	958			43.0%
▼ tstep	958	10000	957.999us	43.0%
► tstep (tstep.c:47)	1			0.04%
► tstep (tstep.c:48)	62			2.78%
► tstep (tstep.c:49)	46			2.06%
► tstep (tstep.c:50)	46			2.06%
► tstep (tstep.c:51)	48			2.15%
► tstep (tstep.c:58)	101			4.53%
► tstep (tstep.c:59)	135			6.06%
► tstep (tstep.c:60)	120			5.39%
► tstep (tstep.c:61)	126			5.66%
► tstep (tstep.c:66)	3			0.13%
► tstep (tstep.c:67)	108			4.85%
► tstep (tstep.c:68)	63			2.83%
► tstep (tstep.c:69)	43			1.93%
► tstep (tstep.c:70)	56			2.51%
► worker.c	2			0.09%

Gprof data in hpctView

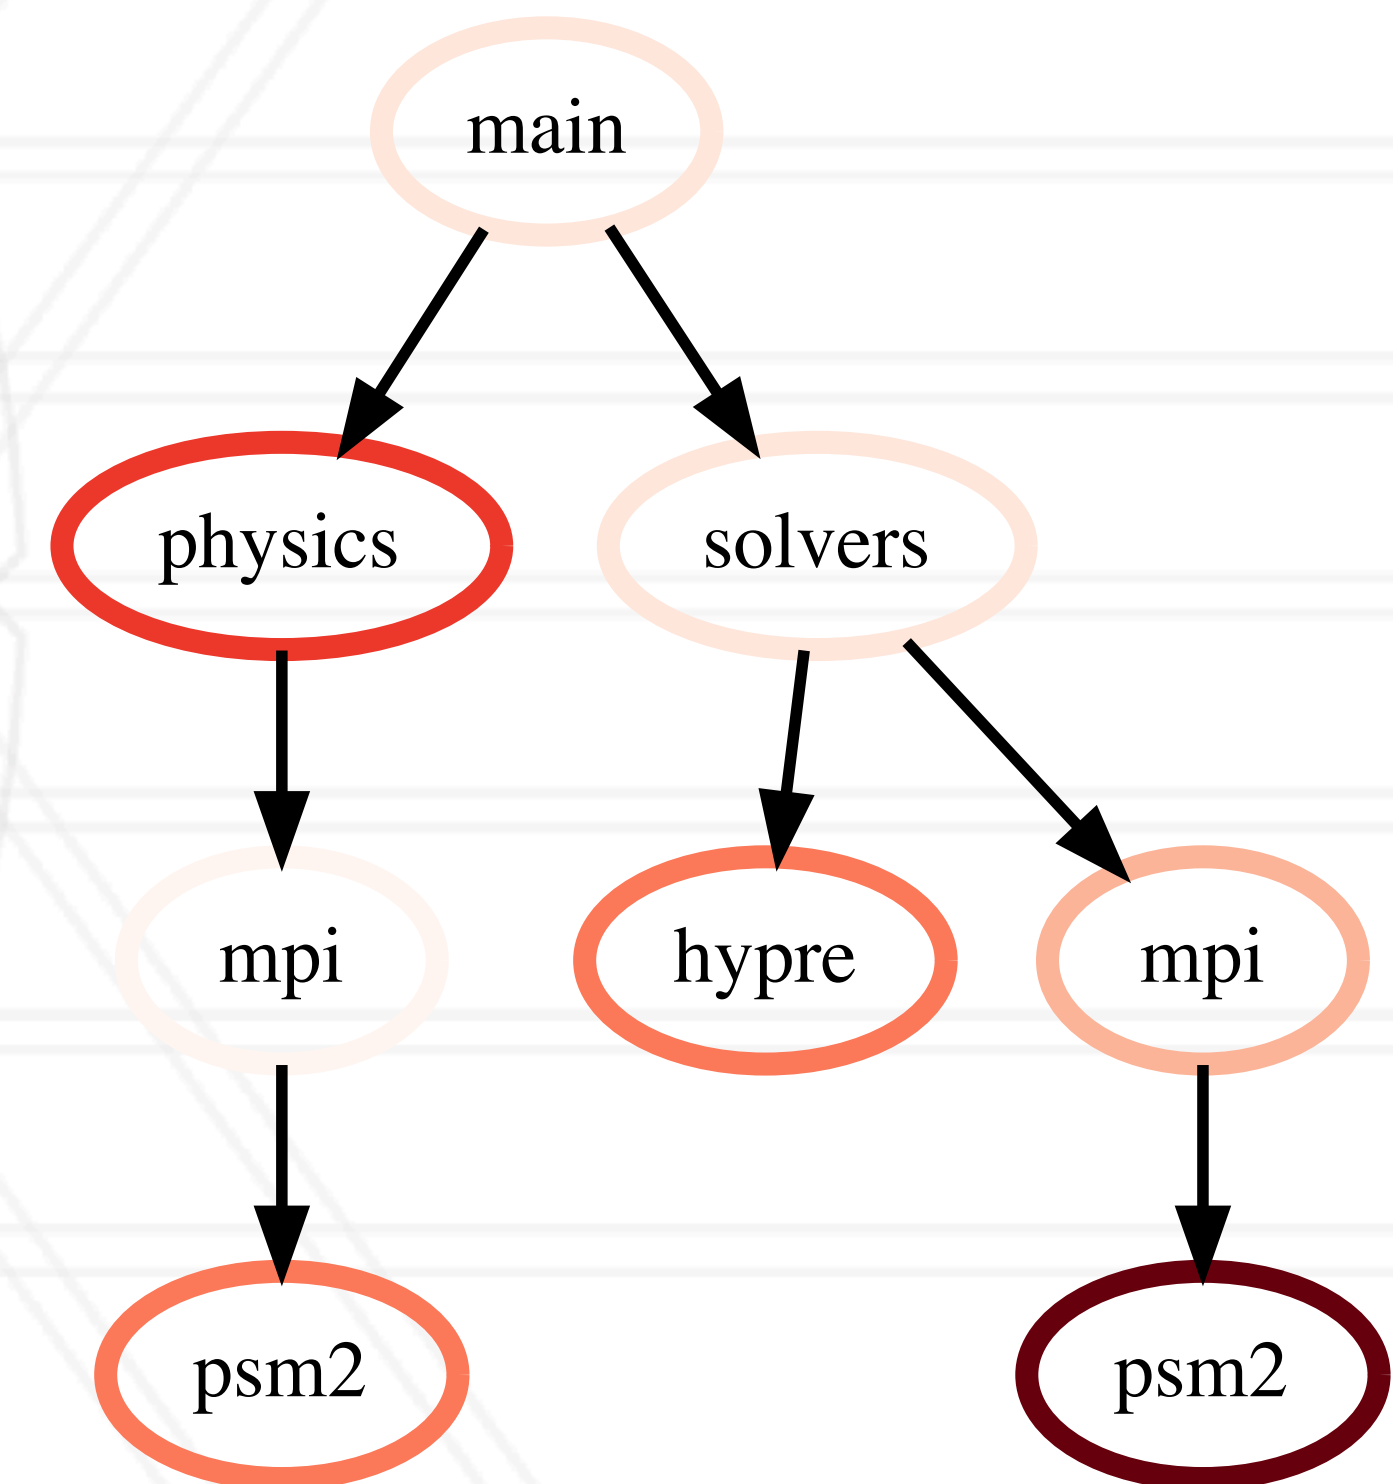


Trace data in Vampir

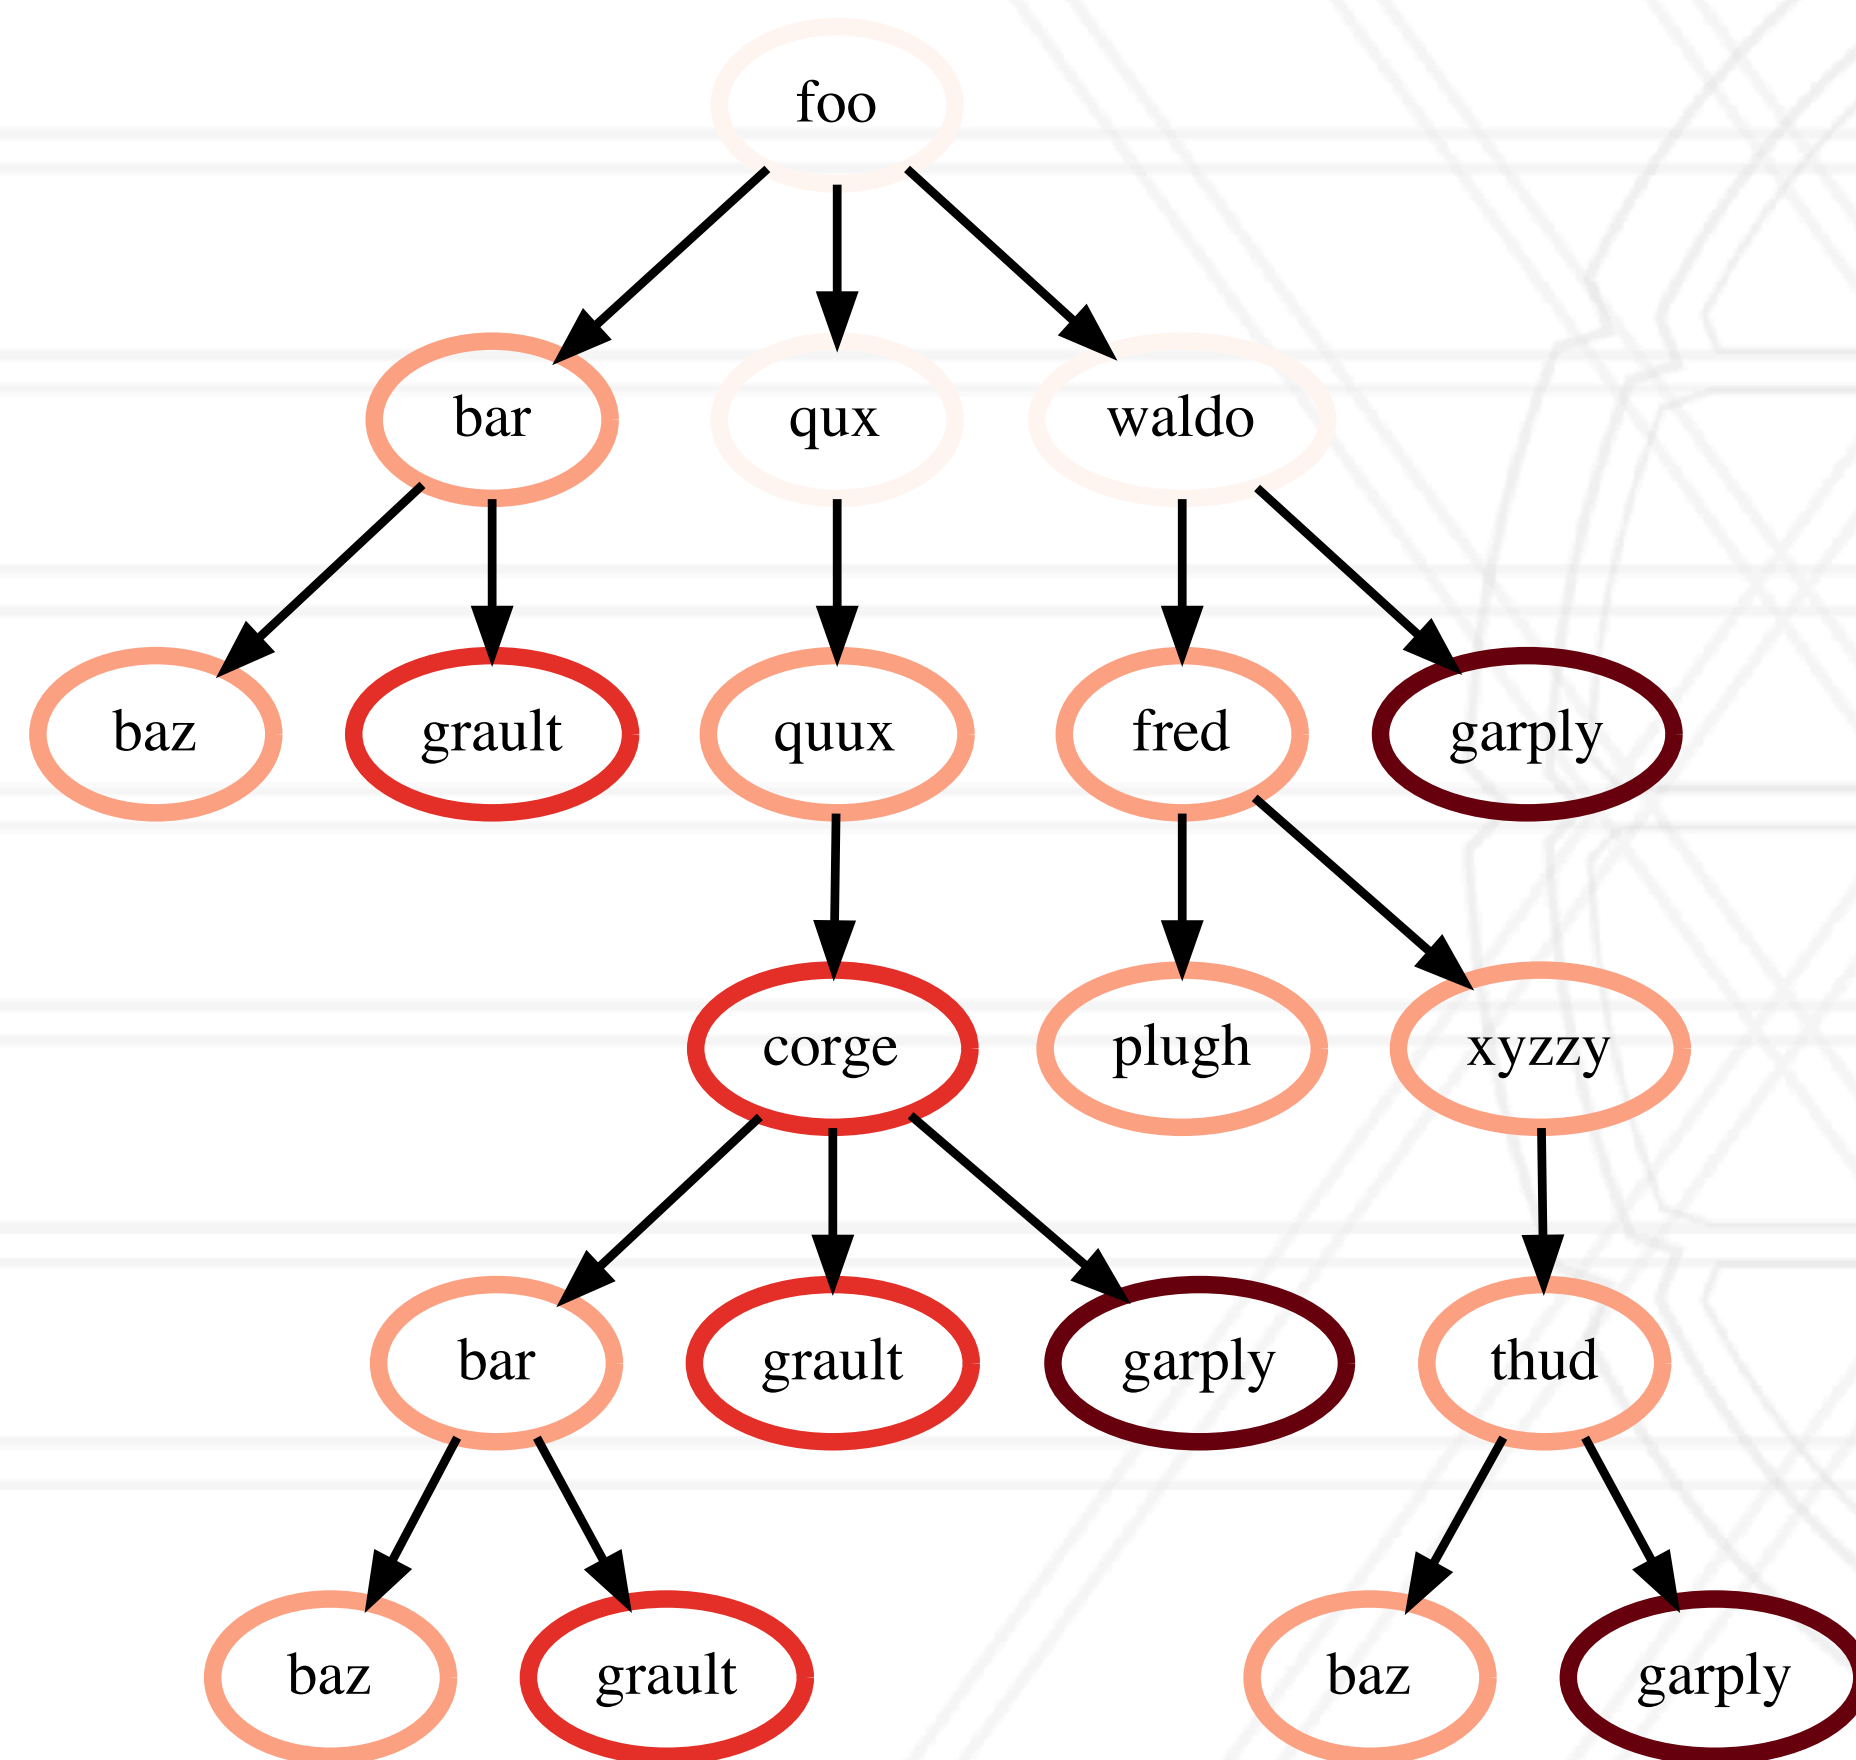


# Attributing performance (time) to different calling contexts

- Requires reasonable understanding of code structure
- Sophisticated profilers can attribute time to calling contexts: where was a function called from?
- Many measurement tools:
  - HPCToolkit, Caliper, Score-P, TAU, Gprof, Callgrind, perf, Timemory
  - cProfile, pyinstrument

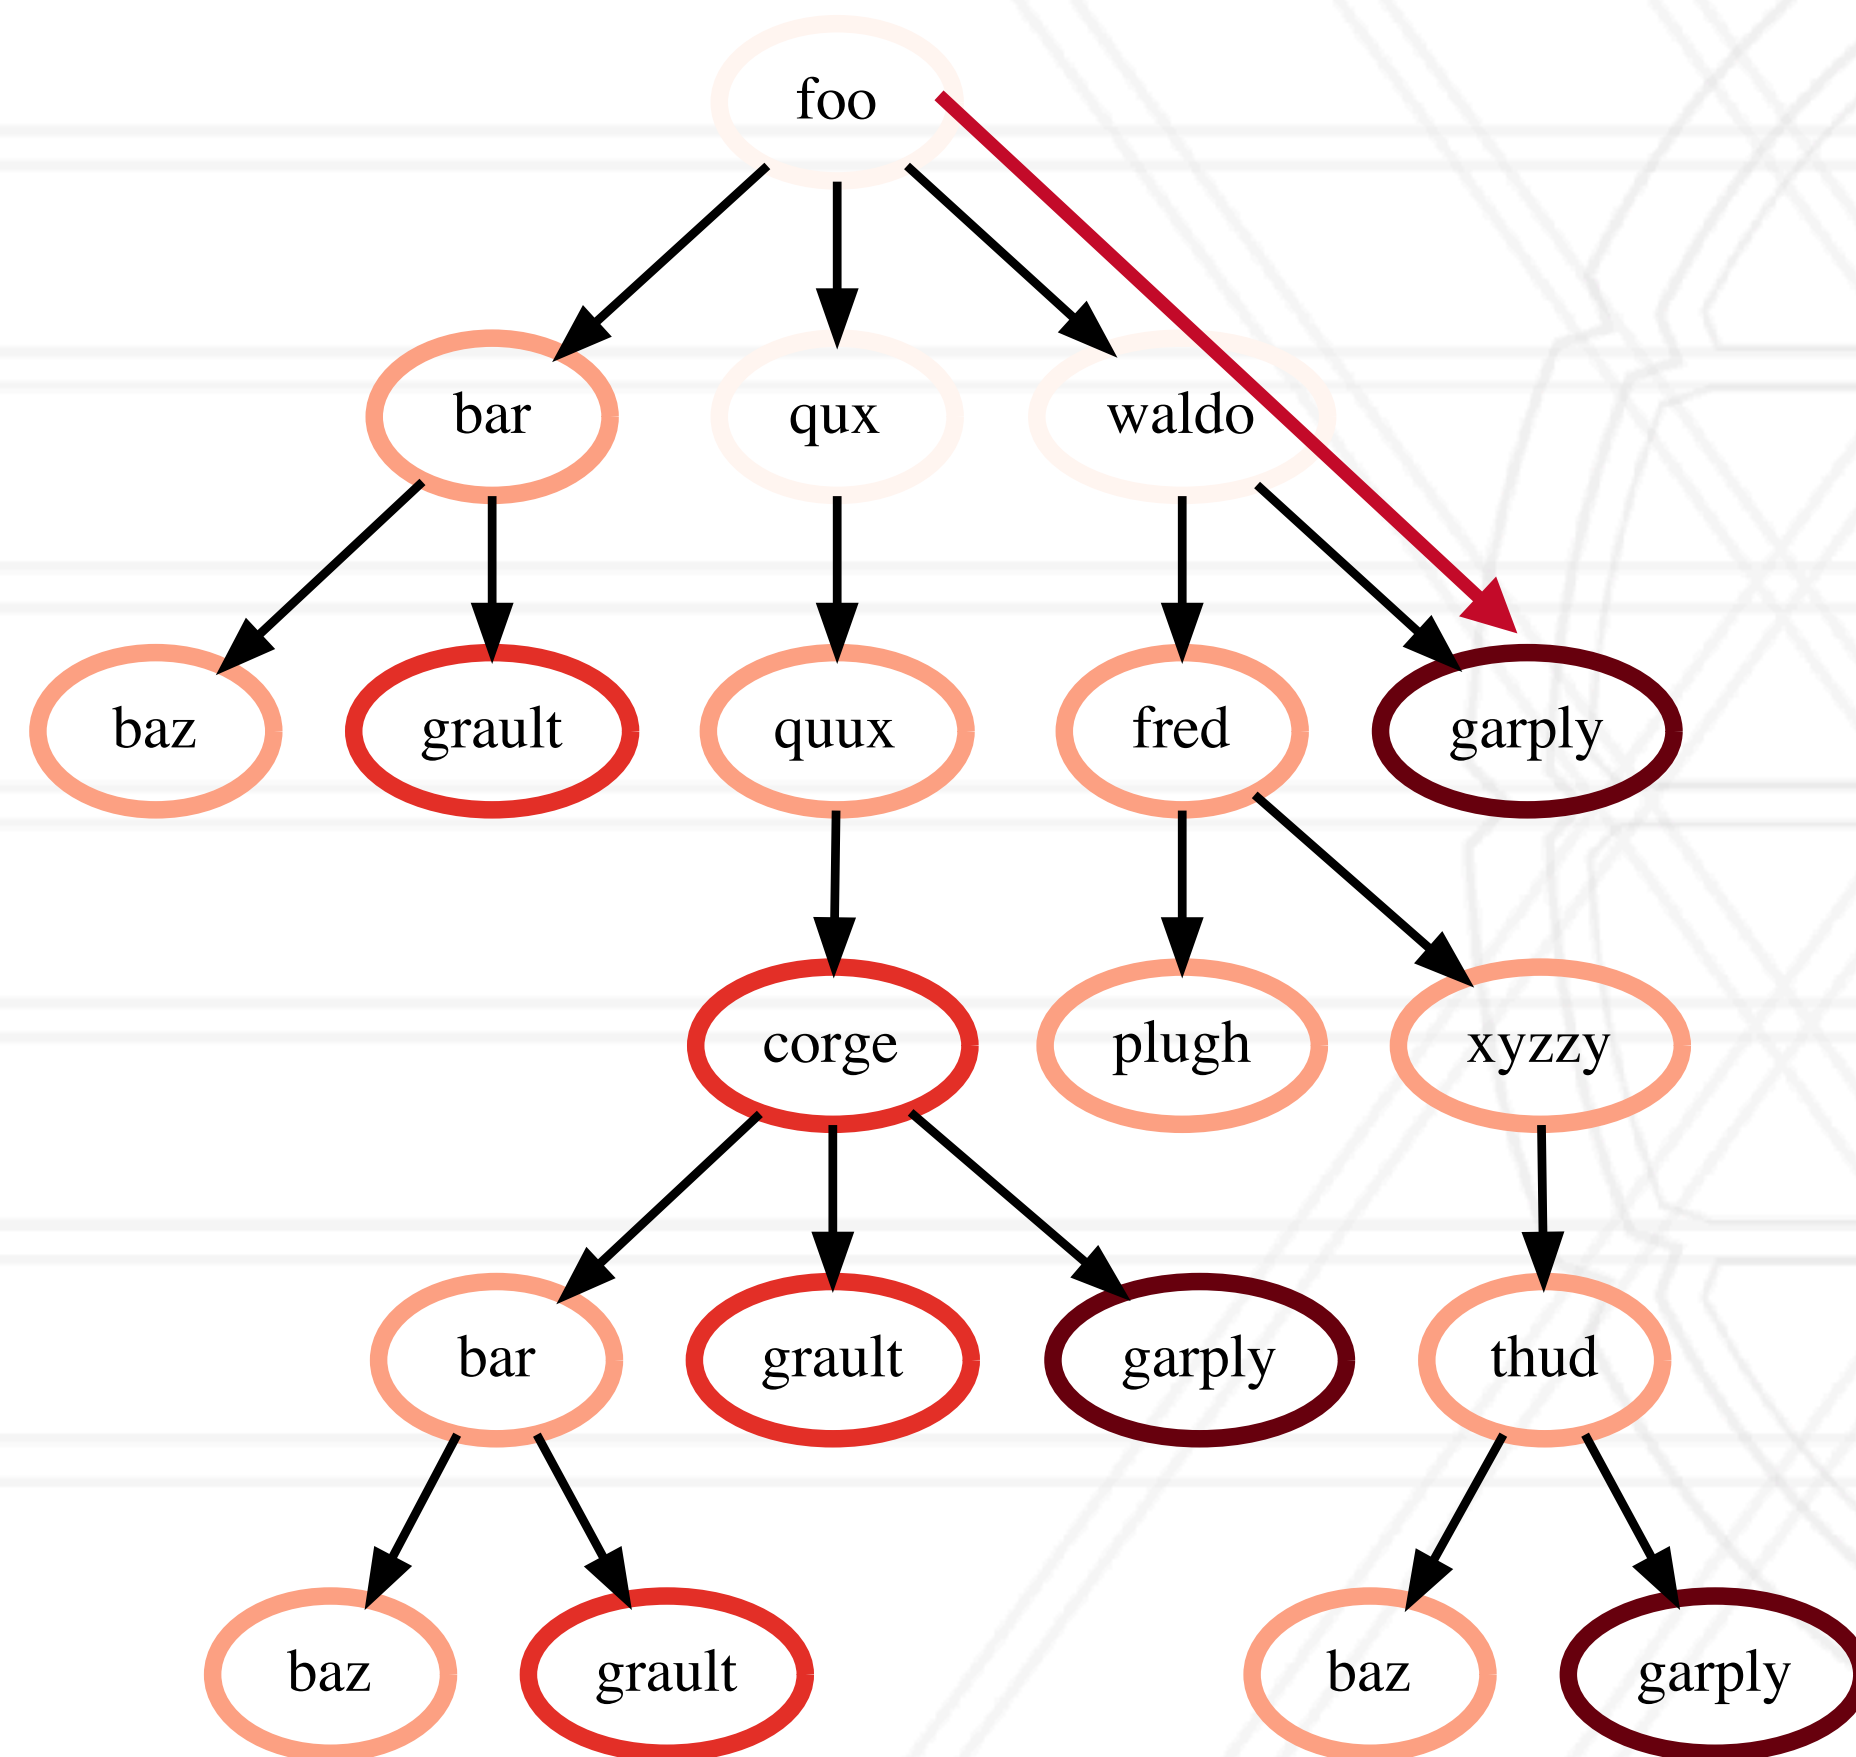


# Calling context trees, call graphs, ...



Calling context tree (CCT)


# Calling context trees, call graphs, ...



Calling context tree (CCT)



# Calling context trees, call graphs, ...




Calling context tree (CCT)

## Contextual information

File  
Line number  
Function name  
Callpath  
Load module  
Process ID  
Thread ID

# Calling context trees, call graphs, ...



Calling context tree (CCT)

## Contextual information

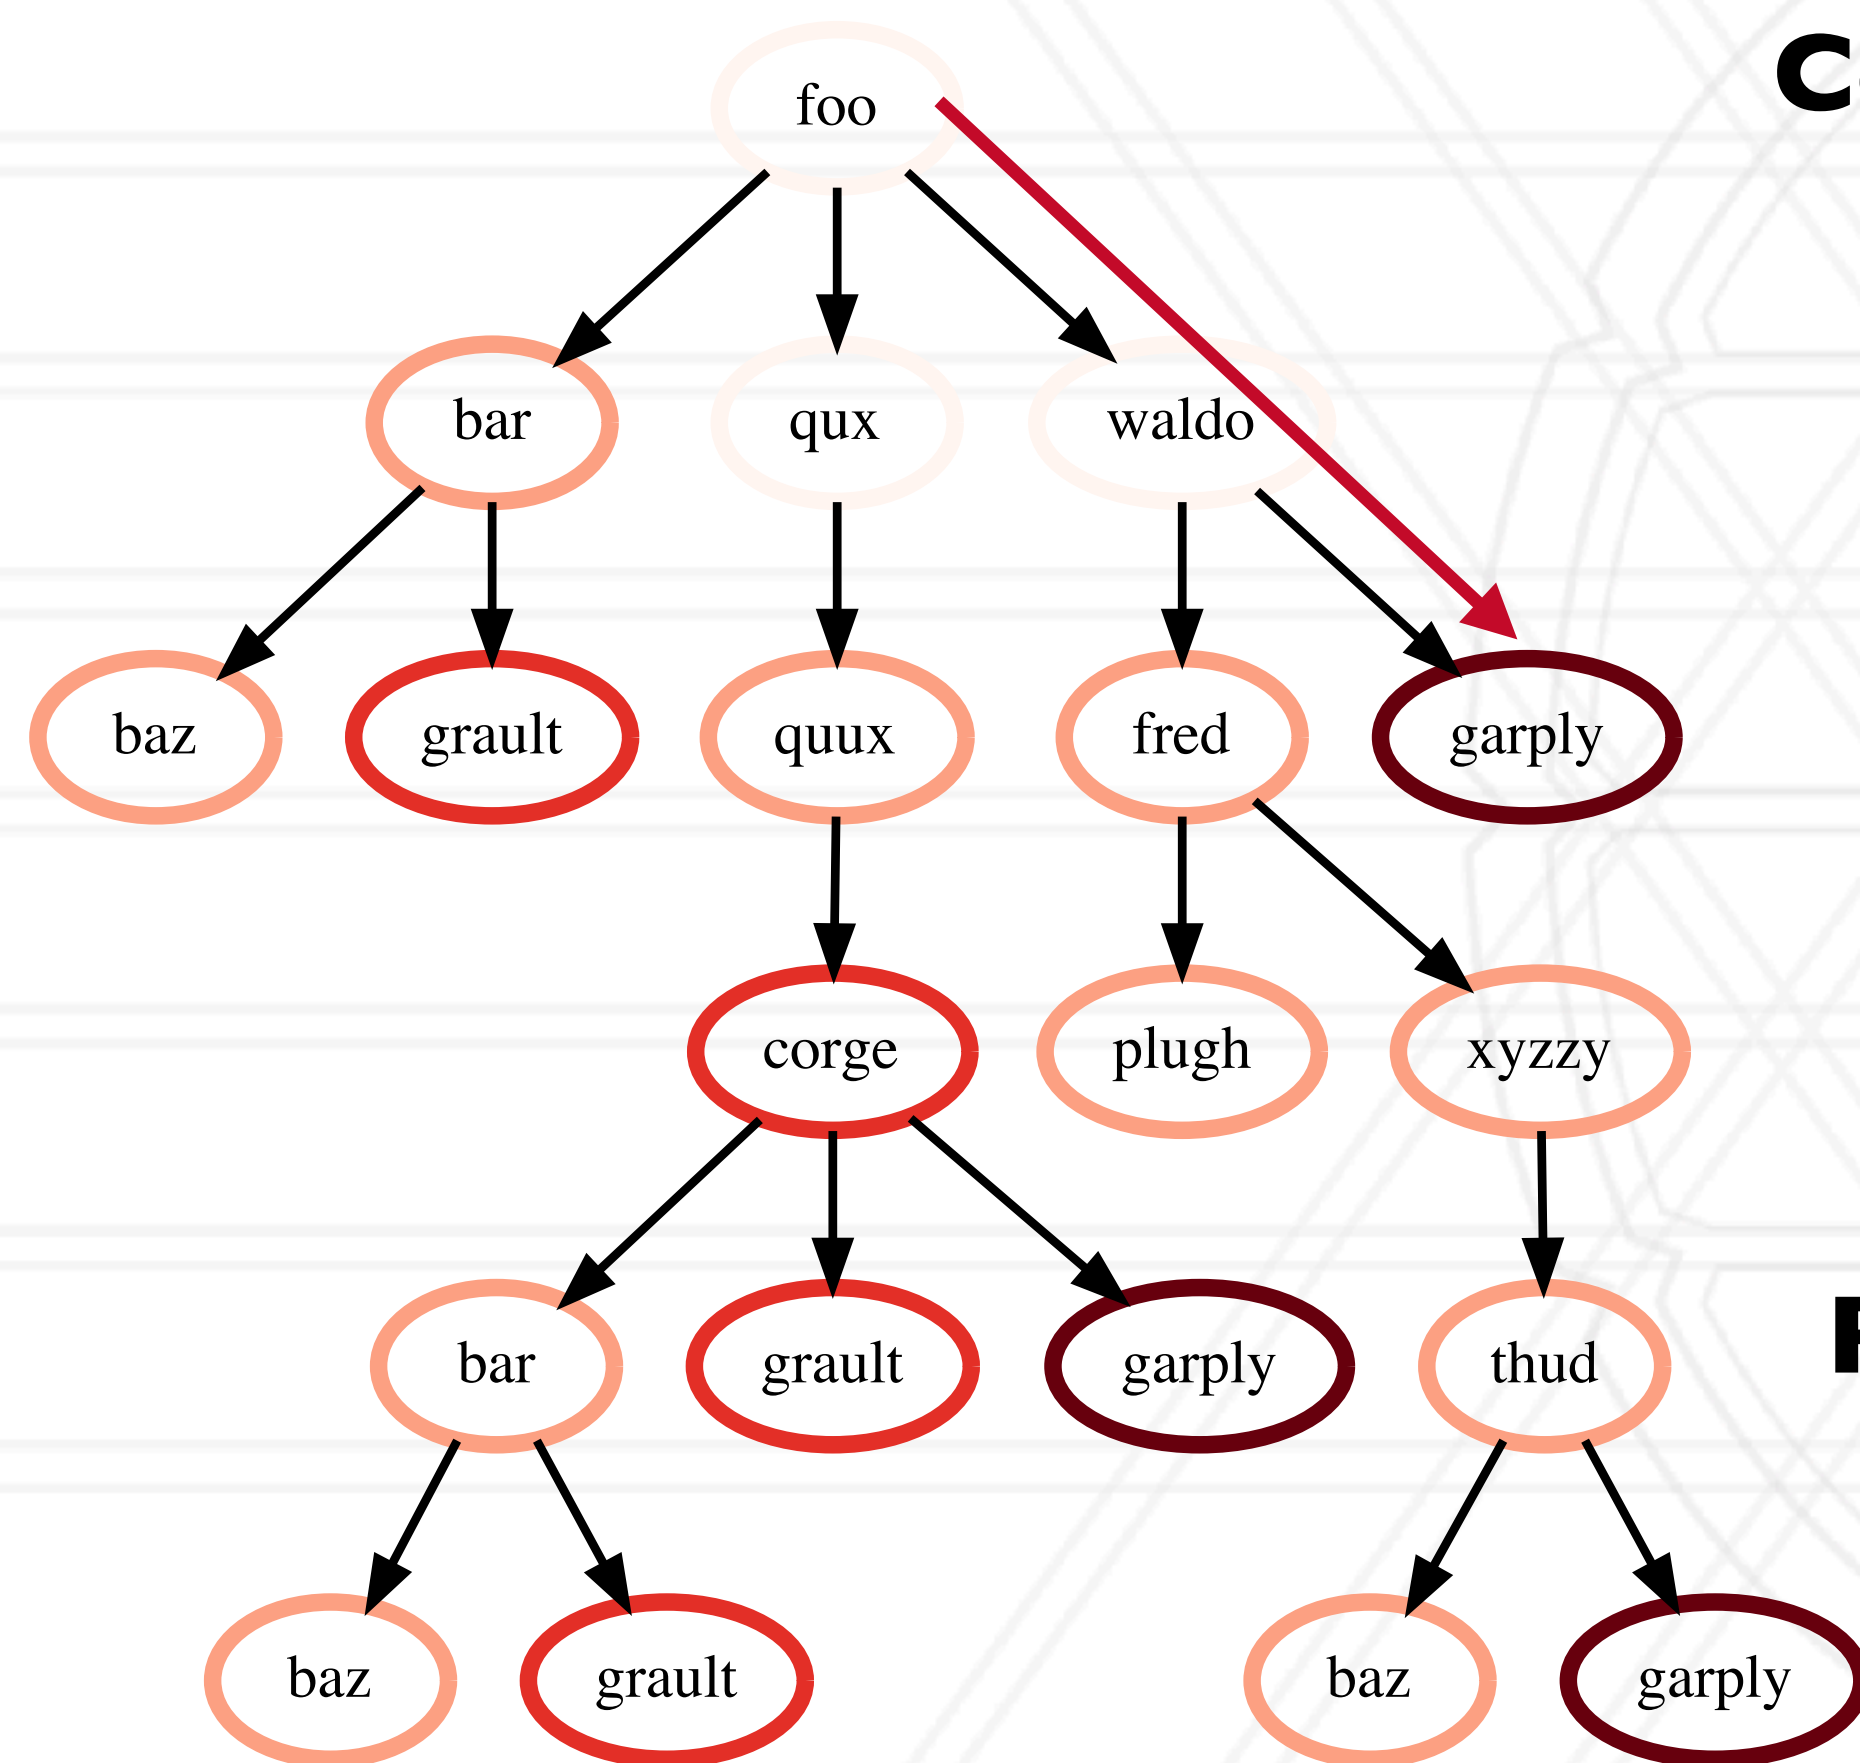
File  
Line number  
Function name  
Callpath  
Load module  
Process ID  
Thread ID

## Performance Metrics

Time  
Flops  
Cache misses



# Calling context trees, call graphs, ...



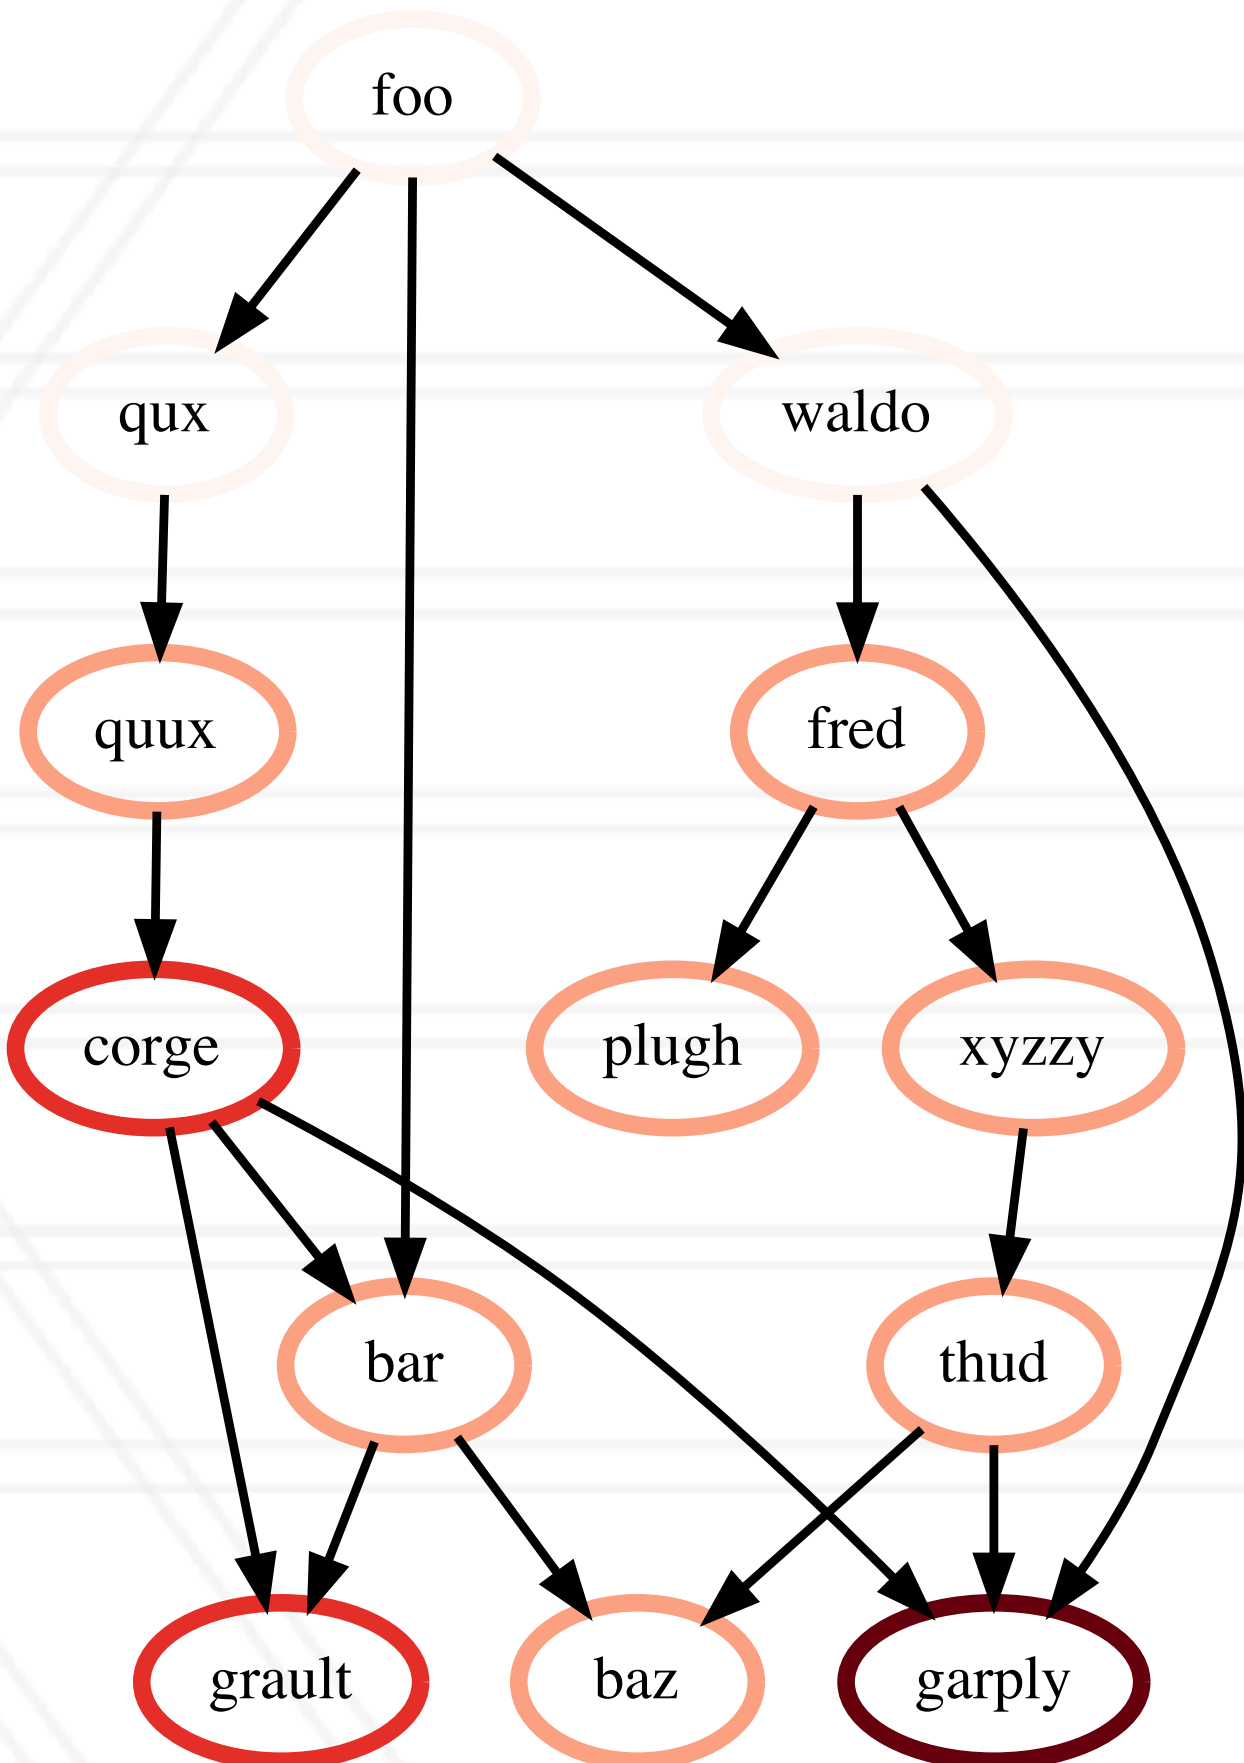
Calling context tree (CCT)

## Contextual information

File  
Line number  
Function name  
Callpath  
Load module  
Process ID  
Thread ID

## Performance Metrics

Time  
Flops  
Cache misses




Call graph



# Limitations of current *analysis* tools

- Support their own unique format(s)
- Limited support for saving or automating analysis
- Most tools only support viewing one dataset at a time
- Lack capabilities to sub-select and focus on specific parts



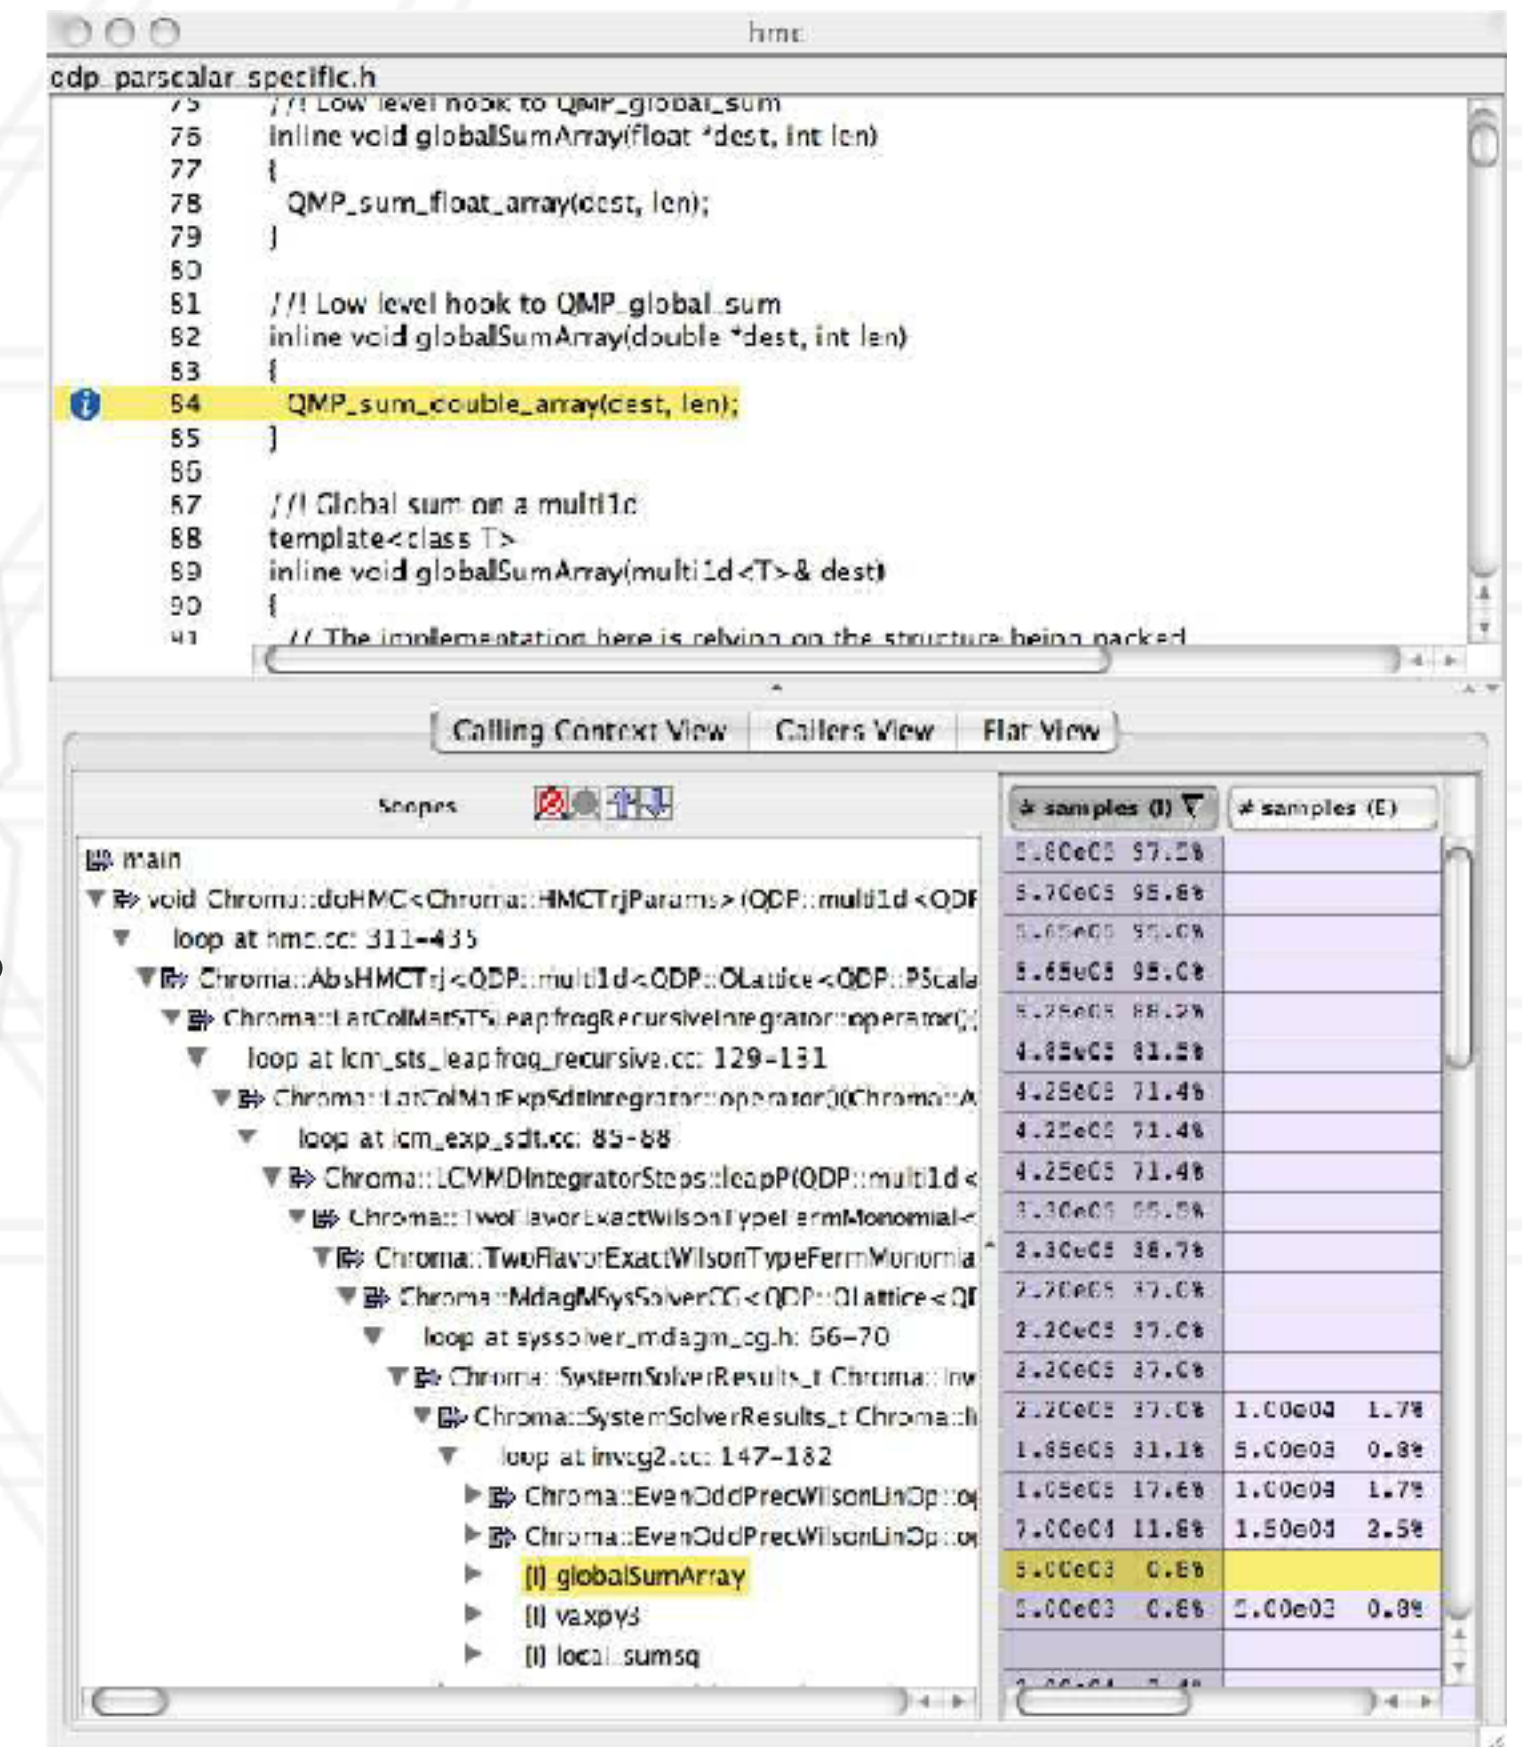
hpcviewer's GUI



# Limitations of current *analysis* tools

- Support their own unique format(s)
- Limited support for saving or automating analysis
- Most tools only support viewing one dataset at a time
- Lack capabilities to sub-select and focus on specific parts

Do not enable programmatic analysis of the data by the end user



hpcviewer's GUI



# The main idea behind Hatchet

---

- A Python-based library to enable programmatic analysis
- Creates an in-memory representation of the graph
- Leverage pandas which supports multi-dimensional tabular datasets
  - Use graph as structured index to index pandas dataframes
- A set of operators to sub-select and/or aggregate profile data
- A set of operators to compare multiple datasets

# Pandas and dataframes

---



# Pandas and dataframes

---

- Pandas is an open-source Python library for data analysis

# Pandas and dataframes

- Pandas is an open-source Python library for data analysis
- Dataframe: two-dimensional tabular data structure
  - Supports many operations borrowed from SQL databases

Columns

	node	name	time (inc)	time
0	{'name': 'main'}	main	200.0	10.0
1	{'name': 'physics'}	physics	60.0	40.0
2	{'name': 'mpi'}	mpi	20.0	5.0
3	{'name': 'psm2'}	psm2	15.0	30.0
4	{'name': 'solvers'}	solvers	100.0	10.0
5	{'name': 'hybre'}	hybre	65.0	30.0
6	{'name': 'mpi'}	mpi	35.0	20.0
7	{'name': 'psm2'}	psm2	25.0	60.0

Rows



# Pandas and dataframes

- Pandas is an open-source Python library for data analysis
- Dataframe: two-dimensional tabular data structure
  - Supports many operations borrowed from SQL databases

Index		Columns			
		node	name	time (inc)	time
Rows	0	{'name': 'main'}	main	200.0	10.0
	1	{'name': 'physics'}	physics	60.0	40.0
	2	{'name': 'mpi'}	mpi	20.0	5.0
	3	{'name': 'psm2'}	psm2	15.0	30.0
	4	{'name': 'solvers'}	solvers	100.0	10.0
	5	{'name': 'hybre'}	hybre	65.0	30.0
	6	{'name': 'mpi'}	mpi	35.0	20.0
	7	{'name': 'psm2'}	psm2	25.0	60.0

# Pandas and dataframes

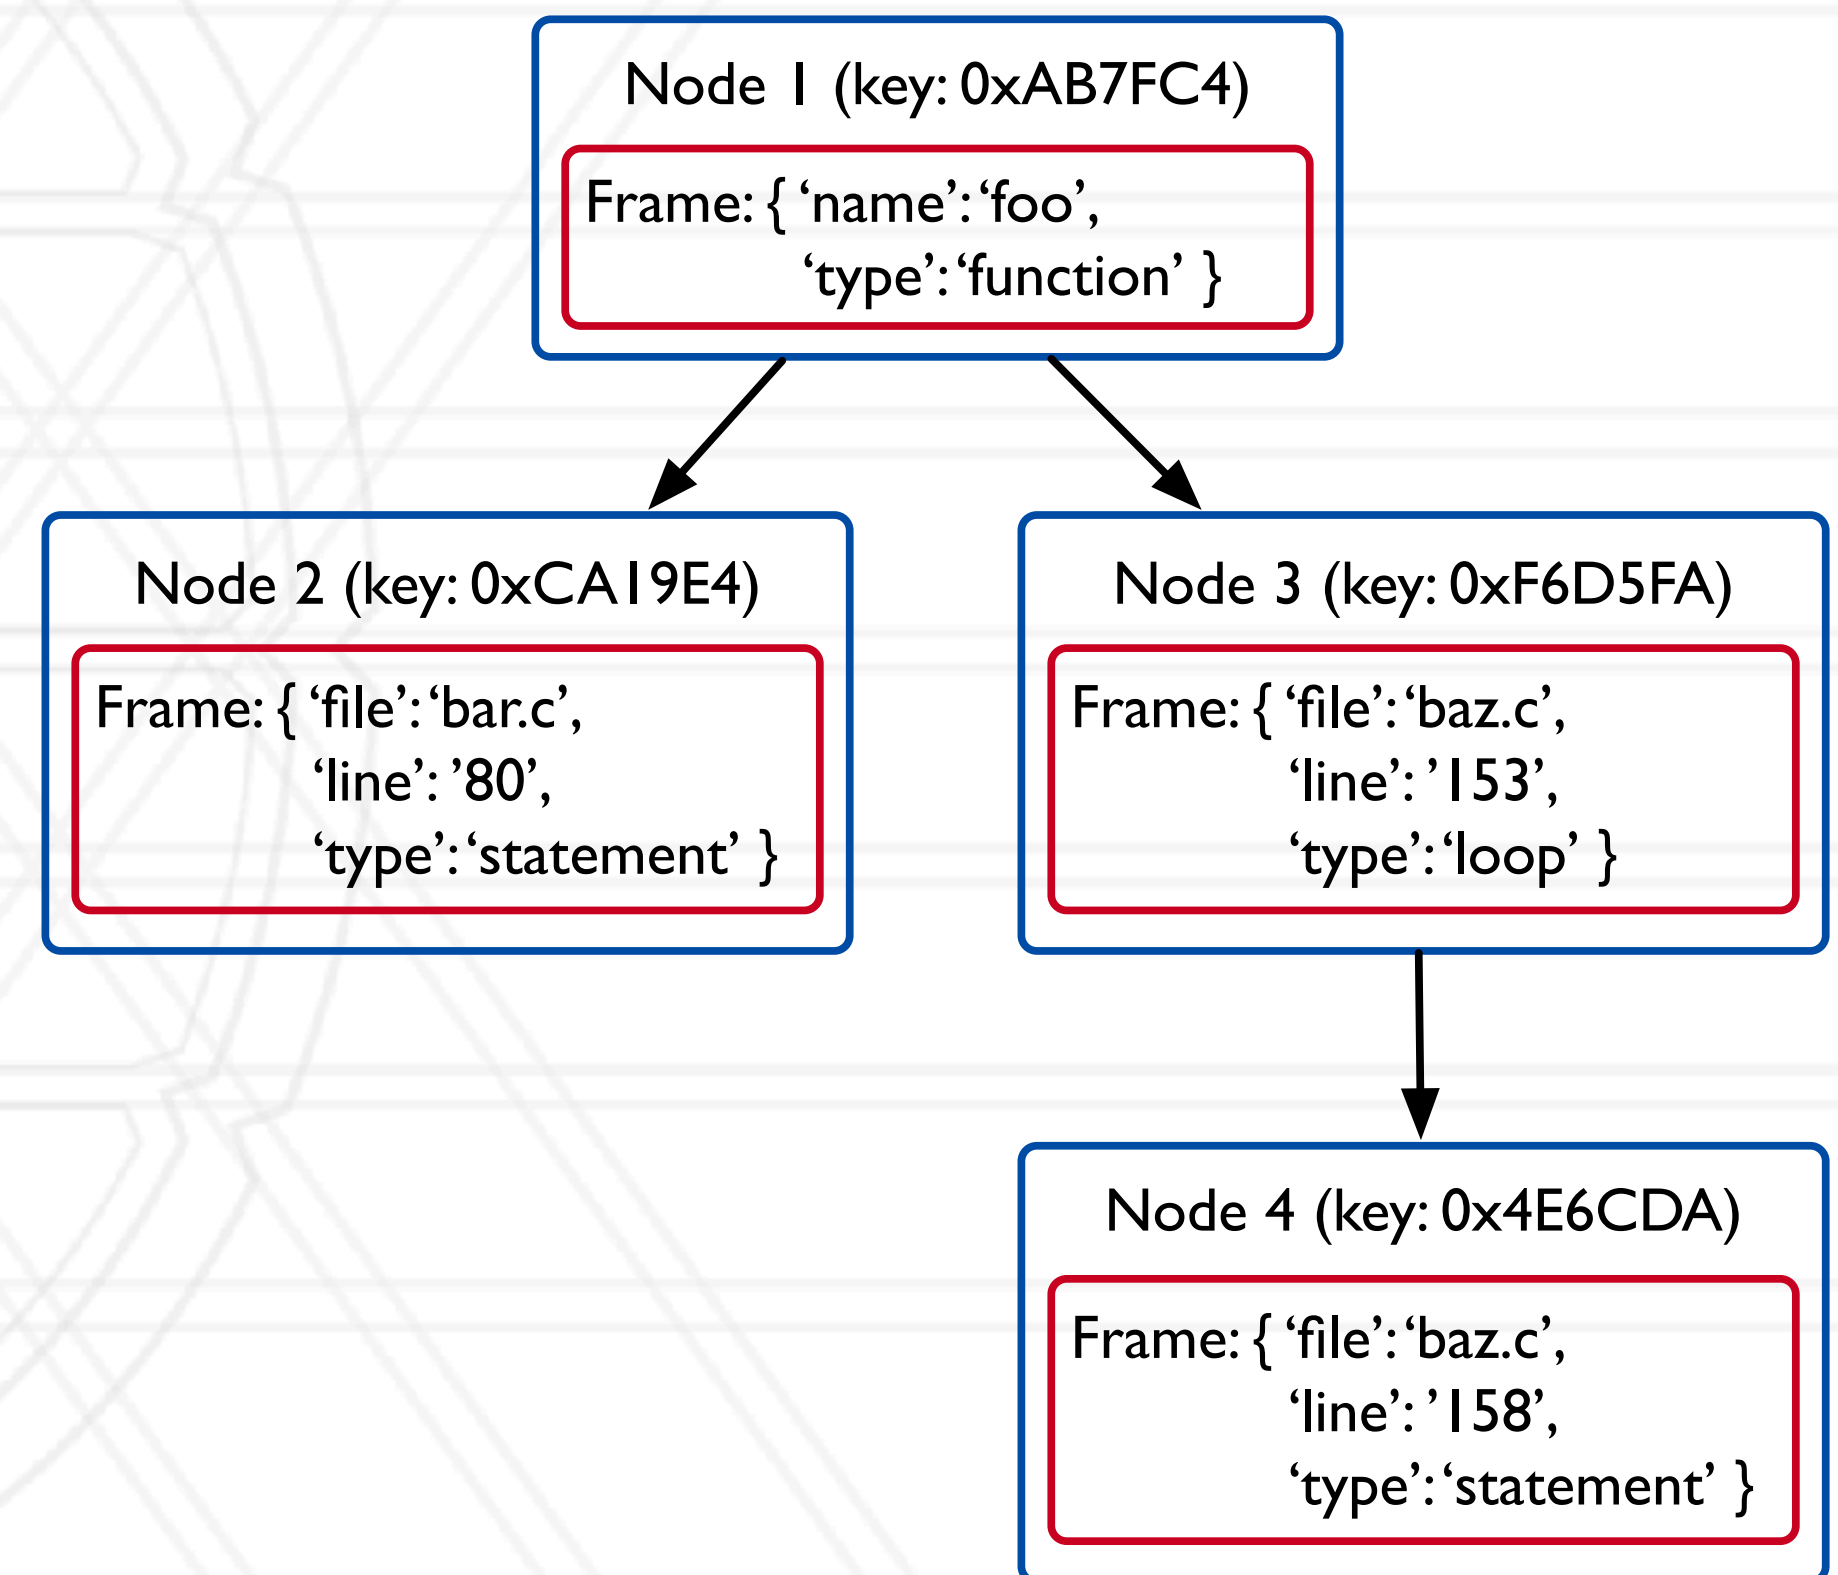
- Pandas is an open-source Python library for data analysis
- Dataframe: two-dimensional tabular data structure
  - Supports many operations borrowed from SQL databases
- MultiIndex enables working with high-dimensional data in a 2D data structure

Index		Columns			
		node	name	time (inc)	time
Rows	0	{'name': 'main'}	main	200.0	10.0
	1	{'name': 'physics'}	physics	60.0	40.0
	2	{'name': 'mpi'}	mpi	20.0	5.0
	3	{'name': 'psm2'}	psm2	15.0	30.0
	4	{'name': 'solvers'}	solvers	100.0	10.0
	5	{'name': 'hybre'}	hybre	65.0	30.0
	6	{'name': 'mpi'}	mpi	35.0	20.0
	7	{'name': 'psm2'}	psm2	25.0	60.0



# Canonical data model: *a structured index*

- Structured index is basically an in-memory graph
- Each node is assigned a unique key, which enables using the nodes as the index in the dataframe
- Each node has a *Frame* that describes the code it represents
  - a set of key/value pairs
- Frames don't have a rigid schema
- *Nodes* define the structure and connectivity



# Central data structure: a *GraphFrame*


---

- Consists of a structured index graph object and a pandas dataframe
- Graph stores caller-callee relationships
- Dataframe stores all numerical and categorical data




# Central data structure: a *GraphFrame*

- Consists of a structured index graph object and a pandas dataframe
- Graph stores caller-callee relationships
- Dataframe stores all numerical and categorical data



# Central data structure: a *GraphFrame*

- Consists of a structured index graph object and a pandas dataframe
- Graph stores caller-callee relationships
- Dataframe stores all numerical and categorical data



	name	nid	node	time	time (inc)
node					
	main	0	main	40.0	200.0
	physics	1	physics	40.0	60.0
	mpi	2	mpi	5.0	20.0
	psm2	3	psm2	15.0	15.0
	solvers	4	solvers	0.0	100.0
	hypre	5	hypre	65.0	65.0
	mpi	6	mpi	10.0	35.0
	psm2	7	psm2	25.0	25.0



# Use of MultiIndex

- When metrics are per MPI process, thread etc....

		time (inc)	time	nid	rank	file	line	module
node	rank							
{ 'type': 'function', 'name': '<program root>' }	0	999238.0	0.0	2	0	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
	1	999390.0	0.0	2	1	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
{ 'type': 'function', 'name': '27:MPI_Init' }	0	0.0	0.0	6	0	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
	1	0.0	0.0	6	1	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
{ 'type': 'function', 'name': 'main' }	0	999238.0	0.0	4	0	./src/cpi.c	19	cpi
	1	999390.0	0.0	4	1	./src/cpi.c	19	cpi

# Use of MultiIndex

- When metrics are per MPI process, thread etc....

Index

		time (inc)	time	nid	rank	file	line	module
	node rank							
{ 'type': 'function', 'name': '<program root>'} }	0	999238.0	0.0	2	0	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
	1	999390.0	0.0	2	1	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
{ 'type': 'function', 'name': '27:MPI_Init'} }	0	0.0	0.0	6	0	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
	1	0.0	0.0	6	1	<unknown file>	0	/collab/usr/global/tools/hpctoolkit/chaos_5_x8...
{ 'type': 'function', 'name': 'main'} }	0	999238.0	0.0	4	0	./src/cpi.c	19	cpi
	1	999390.0	0.0	4	1	./src/cpi.c	19	cpi



# *Immutable* semantics for graph nodes

---

- Having direct references to graph nodes in the dataframe is risky
  - In particular when graph nodes are shared by multiple graphframes
- Any operation that modifies graph nodes in place creates a new GraphFrame and a new graph index
- Implemented using copy-on-write semantics

# Reading in an input dataset

---

## Installing hatchet



```
$ pip install hatchet
```

- HPCToolkit, Caliper, gprof
- Pyinstrument, cprofile
- String literal
- In progress: Timemory, TAU, cube

**<https://github.com/hatchet/hatchet>**



# Reading in an input dataset

## Installing hatchet



```
$ pip install hatchet
```

<https://github.com/hatchet/hatchet>

- HPCToolkit, Caliper, gprof
- Pyinstrument, cprofile
- String literal
- In progress: Timemory, TAU, cube

```
import hatchet as ht

if __name__ == "__main__":
    dirname = "hpctoolkit-database"
    gf = ht.GraphFrame.from_hpctoolkit(dirname)
```

# Reading in an input dataset

## Installing hatchet



```
$ pip install hatchet
```

<https://github.com/hatchet/hatchet>

- HPCToolkit, Caliper, gprof
- Pyinstrument, cprofile
- String literal
- In progress: Timemory, TAU, cube

Contribute a reader  
to hatchet!

```
import hatchet as ht

if __name__ == "__main__":
    dirname = "hpctoolkit-database"
    gf = ht.GraphFrame.from_hpctoolkit(dirname)
```





Abhinav Bhatele

Parallel Software and Systems Group

5218 Brendan Iribe Center (IRB) / College Park, MD 20742

phone: 301.405.4507 / e-mail: [bhatele@cs.umd.edu](mailto:bhatele@cs.umd.edu)



UNIVERSITY OF  
MARYLAND

